


Introduction to Machine Learning

ML-Basics In a Nutshell



Learning goals


- Understand fundamental goal of supervised machine learning
- Know concepts of task, model, parameter, learner, loss function, and empirical risk minimization

WHAT IS ML?


“A computer program is said to learn from experience E with respect to some task T and some performance measure P , if its performance on T , as measured by P , improves with experience E .”

Tom Mitchell, Carnegie Mellon University, 1998

⇒ 99 % of this lecture is about **supervised learning**:



Training



Prediction




TASKS


- Supervised tasks are labeled data situations where the goal is to learn the functional relationship between inputs (features) and output (target)
- We distinguish between **regression** and **classification** tasks, depending on whether the target is **numerical** or **categorical**



Regression: Target is **numerical**, e.g., predict days a patient has to stay in hospital





Classification: Target is **categorical**, e.g., predict one of two risk categories for a life insurance customer



MODELS AND PARAMETERS

- A model is a function that maps features to predicted targets




- For finding the model that describes the relation between features and target best, one needs to restrict the set of all possible functions
- This restricted set of functions is called **hypothesis space**. E.g., one could consider only simple linear functions as hypothesis space
- Functions are fully determined by parameters. E.g., in the case of linear functions, $y = \theta_0 + \theta_1 x$, the parameters θ_0 (intercept) and θ_1 (slope) determine the relationship between y and x
- Finding the optimal model means finding the optimal set of parameters

LEARNER


- Learns automatically the relation between features and target – given a set of training data
- Learner picks the best element of the **hypothesis space**, i.e., the function that fits the training data best



Regression:




Classification:



LEARNER

- Learner uses labeled training data to learn a model f . This model is applied to new data for predicting the target variable



Train Set

y	x_1	x_2
2200	4	4300
1800	12	2700
1920	15	3100

Learner

New Features

x_1	x_2
6	3300
5	3100

Model
 f

Prediction of
Target Variable

\hat{y}
2050
2200


LOSS AND RISK MINIMIZATION

- Loss: Measured pointwise for each observation, e.g., L_2 -loss


$$L(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2$$

- Risk: Measured for entire model. Sums up pointwise losses.


$$\mathcal{R}_{\text{emp}}(f) = \sum_{i=1}^n L(y^{(i)}, f(\mathbf{x}^{(i)}))$$



Squared **loss** of one **observation**.



Empirical **risk** of entire **model**



EMPIRICAL RISK MINIMIZATION

- The risk surface visualizes the empirical risk for all possible parameter values of the parameter vector θ
- Minimizing the empirical risk is usually done by numerical optimization

$$\hat{\theta} = \arg \min_{\theta \in \Theta} \mathcal{R}_{\text{emp}}(\theta).$$

